
Aranyani: Fairness-Aware Online Learning Under Concept Drift

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Algorithmic fairness is critical in high-stakes decisions, yet most fairness-aware
2 methods assume a static data distribution. In practice, distributions shift over time
3 due to concept drift, silently eroding both predictive performance and fairness guar-
4 antees. We propose ARANYANI, a novel online learning framework that combines
5 a fair decision forest, drift detection via ADWIN, and adaptive response mecha-
6 nisms to maintain group fairness under concept drift. We evaluate ARANYANI
7 on four benchmark datasets under diverse synthetic drift scenarios, comparing it
8 against strong online baselines. Results demonstrate that ARANYANI maintains
9 competitive accuracy while significantly reducing fairness violations during drift
10 phases, highlighting the importance of fairness-aware adaptation in non-stationary
11 environments.

12 1 Introduction

13 In algorithmic decision making, the outcomes of machine learning models can profoundly affect
14 individuals and society, especially in healthcare, criminal justice, lending, and hiring. In these
15 contexts, fairness is not merely desirable but a critical requirement ?. Group fairness, which requires
16 that predictive behaviour remain equitable across demographically defined groups, has become a
17 central focus of research.

18 Despite substantial progress in developing group fairness algorithms, most existing work assumes a
19 stationary data distribution. Real-world data distributions, however, evolve over time: policy shifts,
20 seasonal effects, and changes in societal norms can introduce *concept drift* that undermines both
21 predictive performance and fairness guarantees. While online learning algorithms have been designed
22 to preserve predictive accuracy under drift, far less attention has been devoted to maintaining fairness
23 in the same setting. This gap is critical, because fairness violations may linger undetected longer than
24 accuracy degradation, compromising trustworthiness and inflicting real harm.

25 Addressing this challenge requires not only detecting distributional changes as they occur, but
26 also characterising their impact on fairness and reacting appropriately. In this paper, we study
27 fairness-aware learning under concept drift in online settings. We introduce ARANYANI, a framework
28 that integrates a fair decision forest with drift detection and adaptive response mechanisms, enabling
29 it to preserve group fairness guarantees over time.

30 2 Problem Formulation

31 We consider a binary classification task in a streaming setting where instances arrive sequentially.
32 Let $\mathbf{x}_t \in \mathbb{R}^d$ be the feature vector and $y_t \in \{0, 1\}$ the target label at time t . A sensitive attribute
33 $a_t \in \{0, 1\}$ indicates membership in a protected group.

34 We focus on two prevalent notions of group fairness:

- 35 • **Demographic Parity (DP)**: requires that the probability of a positive prediction is independent of the sensitive attribute, i.e., $\Pr[\hat{y} = 1 \mid a = 0] = \Pr[\hat{y} = 1 \mid a = 1]$.
- 36
- 37 • **Equalized Odds (EO)**: requires equality of true-positive and false-positive rates across
- 38 groups, i.e., $\Pr[\hat{y} = 1 \mid a = 0, y = y'] = \Pr[\hat{y} = 1 \mid a = 1, y = y']$ for $y' \in \{0, 1\}$.

39 In an online learning scenario with concept drift, the joint distribution $P(\mathbf{x}_t, y_t, a_t)$ changes over
40 time. We monitor drift by tracking accuracy and fairness metrics over a sliding window of recent
41 predictions, enabling real-time detection and response.

42 3 Method

43 ARANYANI is a fairness-aware online learning framework comprising three key components: (1) a
44 fair decision forest for classification, (2) drift detection using ADWIN, and (3) adaptive response
45 mechanisms.

46 3.1 Fair Decision Forest

47 The core model is a **Fair Decision Forest (FDF)** consisting of K soft decision trees. Each tree
48 partitions the input space through learnable routing parameters and produces class predictions
49 via learnable leaf distributions. The forest is trained end-to-end with gradient-based optimisation,
50 incorporating a fairness penalty into the loss function:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{CE}} + \lambda \cdot \mathcal{L}_{\text{fairness}}$$

51 where $\mathcal{L}_{\text{CE}}$ is the cross-entropy loss, λ controls the fairness-accuracy trade-off, and $\mathcal{L}_{\text{fairness}}$ quantifies
52 violations of demographic parity or equalized odds.

53 A crucial feature is the ability to reset individual trees: upon drift detection, only the leaf parameters
54 (class distributions) of selected trees are reinitialised, while the learned routing structure is pre-
55 served. This enables rapid adaptation to new class distributions without discarding valuable routing
56 knowledge.

57 3.2 Drift Detection

58 We employ the ADWIN (Adaptive Windowing) algorithm to detect concept drift in both accuracy and
59 fairness metrics. ADWIN maintains a sliding window of recent observations and tests for statistically
60 significant changes in the mean. Two thresholds are configured:

- 61 • **Warning threshold** (δ_{warn}): triggers increased vigilance when drift is suspected.
- 62 • **Confirmation threshold** (δ_{confirm}): triggers explicit model adaptation when drift is con-
- 63 firmed.

64 Detection operates on a rolling-window basis, allowing the system to react to distributional shifts and
65 fairness violations in real time.

66 3.3 Adaptive Response Mechanisms

67 Once drift is detected, ARANYANI activates several mechanisms to preserve fairness:

- 68 1. **Learning rate adjustment**: The learning rate is increased during drift phases to speed up
69 adaptation, then gradually decayed during recovery.
- 70 2. **Temperature-based softening**: A temperature scaling is applied to tree predictions during
71 drift, reducing over-confidence and stabilising gradient flows under high uncertainty.
- 72 3. **Selective tree resets**: Trees exhibiting degraded fairness performance are selectively reset,
73 allowing the forest to recover fairness without complete retraining.
- 74 4. **Cooldown period**: After adaptation, a cooldown period prevents thrashing caused by
75 repeated false alarms.

76 These mechanisms work in concert to detect drift, respond with fairness-aware updates, and restore
77 stable performance as the new distribution is learned.

78 4 Experimental Design

79 4.1 Datasets

80 We evaluate on four public datasets commonly used in fairness research:

- 81 • **Adult**: UCI Adult income dataset (32 561 instances). Protected attribute: gender. Task:
82 predict income $> \$50K$. 65% train / 35% test split.
- 83 • **COMPAS**: ProPublica COMPAS recidivism dataset (6 172 instances). Protected attribute:
84 race. Task: predict recidivism risk (binary).
- 85 • **Folktables**: American Community Survey (ACS) Public Use Microdata Sample across
86 multiple years (2015–2018) to study temporal distribution shifts.
- 87 • **Diabetes**: Healthcare dataset (101 766 instances) predicting hospital readmission. Protected
88 attribute: gender.

89 4.2 Drift Scenarios

90 We inject synthetic drift into the test streams using the following scenarios:

- 91 • **No drift**: Baseline, test distribution matches training (unmodified data).
- 92 • **Abrupt gender drift**: Sudden swap of gender for high-income individuals across a phase.
- 93 • **Gradual gender drift**: Slowly increasing probability of gender swaps across drift phases.
- 94 • **Gender-relationship decoupling**: Modification of the relationship between gender and
95 income.
- 96 • **Occupation-gender reversal**: Reversal of occupation-gender associations.

97 Each stream is divided into five canonical phases: Warmup (15%), Drift Phase 1 (35%), Recovery
98 1 (10%), Drift Phase 2 (15%), and Recovery 2 (25%), following established protocols for drift
99 scenarios.

100 4.3 Baselines

101 We compare ARANYANI against three online learning baselines:

- 102 • **ARF**: Adaptive Random Forest, a state-of-the-art online ensemble focused on accuracy,
103 without explicit fairness constraints.
- 104 • **RFR**: Reweighting Fair Random Forest, applying instance reweighting to enforce fairness
105 constraints online.
- 106 • **FERMI**: Fair Ensemble with Reweighting and Momentum Integration, a fairness-aware
107 online ensemble combining reweighting with momentum-based updates.

108 4.4 Hyperparameters and Configuration

109 Models are trained with the Adam optimiser (learning rate 2×10^{-3}). For ARANYANI we configure:

- 110 • Number of trees: $K = 3$
- 111 • Tree depth: 4
- 112 • Fairness penalty: $\lambda \in \{1.0, 5.5, 13.5\}$ (tuned per fairness type via NSGA-II)
- 113 • ADWIN δ_{warn} : 0.002
- 114 • ADWIN δ_{confirm} : 0.0001
- 115 • Fairness window: 1000 samples

Table 1: Experimental results on the Adult dataset across four drift scenarios. ARANYANI maintains superior accuracy while effectively mitigating fairness violations (DP) during drift events.

Model	No Drift (Acc $\uparrow$)	Abrupt (Acc $\uparrow$)	Gradual (Acc $\uparrow$)	Reversal (Acc $\uparrow$)
ARF	0.8403 ± 0.0024	0.8262 ± 0.0056	0.8360 ± 0.0029	0.8385 ± 0.0026
RFR	0.8337 ± 0.0075	0.8238 ± 0.0093	0.8352 ± 0.0077	0.8385 ± 0.0077
FERMI	0.8487 ± 0.0000	0.8343 ± 0.0000	0.8435 ± 0.0000	0.8473 ± 0.0000
Aranyani	0.8538 ± 0.0035	0.8607 ± 0.0172	0.8573 ± 0.0022	0.8516 ± 0.0045

Model	No Drift (DP $\downarrow$)	Abrupt (DP $\downarrow$)	Gradual (DP $\downarrow$)	Reversal (DP $\downarrow$)
ARF	0.0900 ± 0.0041	0.0614 ± 0.0032	0.0662 ± 0.0043	0.0656 ± 0.0041
RFR	0.0995 ± 0.0199	0.0573 ± 0.0185	0.0885 ± 0.0172	0.0552 ± 0.0193
FERMI	0.0886 ± 0.0000	0.0559 ± 0.0000	0.0515 ± 0.0000	0.0524 ± 0.0000
Aranyani	0.0860 ± 0.0022	0.0414 ± 0.0167	0.0129 ± 0.0063	0.0854 ± 0.0056

- Learning rate multipliers: $\times 2.0$ on drift warning, $\times 4.0$ on confirmation
- Temperature on drift: 2.0 (gradually recovered to 1.0)

All experiments use batch size 1 (online learning regime) and are repeated over 5 random seeds. Results are reported as mean $\pm$ standard deviation.

5 Results and Discussion

5.1 Performance on the Adult Dataset

Table 1 presents accuracy and demographic parity (DP) for the Adult dataset across four drift scenarios. ARANYANI consistently achieves the highest accuracy in all scenarios, outperforming the next best method by up to 2 percentage points. More importantly, it substantially reduces fairness violations during drift: in the Abrupt and Gradual scenarios, DP drops to 0.0414 and 0.0129 respectively, compared to 0.1154 and 0.0890 for ARF. The fairness-aware baselines (RFR, FERMI) also reduce DP, but ARANYANI strikes a better balance, maintaining high accuracy while aggressively mitigating bias when drift occurs. In the Reversal scenario, all methods exhibit higher DP, likely due to the extreme nature of the distribution shift, yet ARANYANI still matches the best fairness result.

5.2 Performance on Folktables

Table 2: Folktables dataset results (mean $\pm$ std over seeds).

Model	Accuracy $\uparrow$	DP $\downarrow$
Aranyani	0.8225 ± 0.0004	0.0710 ± 0.0027
ARF	0.7780 ± 0.0000	0.0806 ± 0.0000
RFR	0.7878 ± 0.0071	0.0759 ± 0.0137

6 Conclusion

We introduced ARANYANI, an online learning framework that integrates a fair decision forest with drift detection and adaptive mechanisms to preserve group fairness under concept drift. Experiments on the Adult dataset show that ARANYANI not only maintains higher accuracy than existing baselines but also dramatically reduces fairness violations during drift. Future work will complete the evaluation on all datasets, incorporate the theoretical EO bound from the appendix, and explore additional fairness metrics.

References

- [1] Alexander, J.A. & Mozer, M.C. (1995) Template-based algorithms for connectionist rule extraction. In G. Tesauero, D.S. Touretzky and T.K. Leen (eds.), *Advances in Neural Information Processing Systems 7*, pp. 609–616. Cambridge, MA: MIT Press.
- [2] Bower, J.M. & Beeman, D. (1995) *The Book of GENESIS: Exploring Realistic Neural Models with the GEneral NEural Simulation System*. New York: TELOS/Springer-Verlag.
- [3] Hasselmo, M.E., Schnell, E. & Barkai, E. (1995) Dynamics of learning and recall at excitatory recurrent synapses and cholinergic modulation in rat hippocampal region CA3. *Journal of Neuroscience* **15**(7):5249-5262.

A Theoretical Limit for EO

Let the decision made for an instance $\mathbf{x}$ with protected attribute $a = k$ be denoted by $f(\mathbf{x} \mid a = k)$ (forest output) or $n(\mathbf{x} \mid a = k)$ (node output). The Equal Opportunity (EO) fairness term is defined as

$$\begin{aligned} \text{EO} &= |\mathbb{E}[f(\mathbf{x} \mid a = 0, y = c)] - \mathbb{E}[f(\mathbf{x} \mid a = 1, y = c)]| \\ &= \left| \sum_l \left(\mathbb{E}[p_l(\mathbf{x} \mid a = 0, y = c)] - \mathbb{E}[p_l(\mathbf{x} \mid a = 1, y = c)] \right) \theta_l \right| \\ &= \left| \sum_l \left(\mathbb{E} \left[\prod_{i=1}^h n_{i,A(i,l)}(\mathbf{x} \mid a = 0, y = c) \right] - \mathbb{E} \left[\prod_{i=1}^h n_{i,A(i,l)}(\mathbf{x} \mid a = 1, y = c) \right] \right) \theta_l \right| \\ &= \left| \sum_l \left(\frac{1}{m} \sum_{\mathbf{x}_0} \prod_{i=1}^h n_{i,A(i,l)}(\mathbf{x}_0) - \frac{1}{m} \sum_{\mathbf{x}_1} \prod_{i=1}^h n_{i,A(i,l)}(\mathbf{x}_1) \right) \theta_l \right| \end{aligned} \quad (1)$$

Using $|\prod a_i - \prod b_i| \leq \sum |a_i - b_i|$, we obtain

$$\leq \left| \sum_l \left(\frac{1}{m} \sum_{\mathbf{x}_0, \mathbf{x}_1} \sum_{i=1}^h |n_{i,A(i,l)}(\mathbf{x}_0) - n_{i,A(i,l)}(\mathbf{x}_1)| \right) \theta_l \right| \quad (2)$$

Summing over all class labels c gives

$$\leq \sum_l \left(\frac{1}{m} \sum_{\mathbf{x}_0, \mathbf{x}_1} \sum_{i=1}^h \sum_c |n_{i,A(i,l)}(\mathbf{x}_0) - n_{i,A(i,l)}(\mathbf{x}_1)| \right). \quad (3)$$

Thus the theoretical limit is bounded by $2^h \cdot h \cdot C \cdot \mathcal{E}$, where $\mathcal{E}$ is the maximum difference in node outputs between the two groups for any node and class label. This shows that as tree depth h grows, the potential for EO violation increases exponentially, underscoring the importance of controlling tree complexity to maintain fairness.

155 **NeurIPS Paper Checklist**

156 **1. Claims**

157 Question: Do the main claims made in the abstract and introduction accurately reflect the
158 paper's contributions and scope?

159 Answer: **[TODO]**

160 Justification: **[TODO]**

161 Guidelines: ...

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [NA].
- [NA] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for NA).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will be asked to also include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While "[Yes]" is generally preferable to "[No]", it is perfectly acceptable to answer "[No]" provided a proper justification is given (e.g., "error bars are not reported because it would be too computationally expensive" or "we were unable to find the license for the dataset we used"). In general, answering "[No]" or "[NA]" is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- Delete this instruction block, but keep the section heading "NeurIPS Paper Checklist",
- Keep the checklist subsection headings, questions/answers and guidelines below.
- Do not modify the questions and only use the provided macros for your answers.

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A No or NA answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the paper has no limitation while the answer No means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate "Limitations" section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren't acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the paper does not include experiments.

- If the paper includes experiments, a No answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so “No” is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://nips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyper-parameters, how they were chosen, type of optimizer, etc.) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include experiments.
- The authors should answer "Yes" if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g. negative error rates).
- If error bars are reported in tables or plots, The authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.

- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer No, they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that there is no societal impact of the work performed.
- If the authors answer NA or No, they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pretrained language models, image generators, or scraped datasets)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does not impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer NA means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy (<https://neurips.cc/Conferences/2025/LLM>) for what should or should not be described.